

The graph displays throughput (in MB/s) on the y-axis against the number of clients (32, 64, 128, 256, 400) on the x-axis. Ten different configurations are compared, each represented by a colored line with error bars at specific client counts. The configurations are listed in the legend:

- 4.19.0-ipanema-g9ba5ed25b696,Linux,Linux
- 4.19.0-ipanema-g9ba5ed25b696,cfs_wwc scheduling policy,cfs_wwc scheduling policy
- 4.19.0-ipanema-g9ba5ed25b696,cfs_wwc_flat scheduling policy,cfs_wwc_flat scheduling policy
- 4.19.0-ipanema-g9ba5ed25b696,uile scheduling policy,uile scheduling policy
- 4.19.0-ipanema-g9ba5ed25b696,uile_wwc scheduling policy,uile_wwc scheduling policy
- 4.19.0-ipanema-gab29e103e36b,Linux,Linux
- 4.19.0-ipanema-gab29e103e36b,cfs_wwc scheduling policy,cfs_wwc scheduling policy
- 4.19.0-ipanema-gab29e103e36b,cfs_wwc_flat scheduling policy,cfs_wwc_flat scheduling policy
- 4.19.0-ipanema-gab29e103e36b,uile scheduling policy,uile scheduling policy
- 4.19.0-ipanema-gab29e103e36b,uile_wwc scheduling policy,uile_wwc scheduling policy

All configurations show a positive linear trend, with throughput increasing as the number of clients increases. The green line (4.19.0-ipanema-g9ba5ed25b696,cfs_wwc_flat scheduling policy,cfs_wwc_flat scheduling policy) generally shows the highest throughput across most client counts.

