

The graph displays the throughput (ops/sec) on the y-axis against the number of clients on the x-axis. The x-axis has major ticks at 32, 64, 128, 256, and 400. The y-axis ranges from 0 to 1000 ops/sec. There are ten data series, each representing a different combination of hardware configuration and scheduling policy. The series are color-coded and labeled in the legend:

- 4.19.0-ipanema-g9ba5ed25b696, Linux, Linux (Blue)
- 4.19.0-ipanema-g9ba5ed25b696, cfs_wwc scheduling policy, cfs_wwc scheduling policy (Orange)
- 4.19.0-ipanema-g9ba5ed25b696, cfs_wwc_flat scheduling policy, cfs_wwc_flat scheduling policy (Green)
- 4.19.0-ipanema-g9ba5ed25b696, ule scheduling policy, ule scheduling policy (Red)
- 4.19.0-ipanema-g9ba5ed25b696, ule_wwc scheduling policy, ule_wwc scheduling policy (Purple)
- 4.19.0-ipanema-gab29e103e36b, Linux, Linux (Brown)
- 4.19.0-ipanema-gab29e103e36b, cfs_wwc scheduling policy, cfs_wwc scheduling policy (Pink)
- 4.19.0-ipanema-gab29e103e36b, cfs_wwc_flat scheduling policy, cfs_wwc_flat scheduling policy (Grey)
- 4.19.0-ipanema-gab29e103e36b, ule scheduling policy, ule scheduling policy (Yellow)
- 4.19.0-ipanema-gab29e103e36b, ule_wwc scheduling policy, ule_wwc scheduling policy (Cyan)

Key observations from the graph:

- Throughput increases linearly with the number of clients for all configurations.
- The **cfs_wwc_flat** scheduling policy (Green line) consistently achieves the highest throughput across all client counts.
- The **ule_wwc** scheduling policy (Purple and Cyan lines) consistently achieves the lowest throughput.
- The hardware configuration **4.19.0-ipanema-gab29e103e36b** (lines 6-10) generally performs better than **4.19.0-ipanema-g9ba5ed25b696** (lines 1-5).
- At 256 clients, there are visible error bars for the Green, Red, and Purple lines, indicating variability in the measurements.

