

Figure 1: Number of clients vs. scheduling policy. The plot shows the performance of various scheduling policies as the number of clients increases from 32 to 512. The y-axis represents a metric (likely latency or throughput) ranging from 0 to 100. The x-axis represents the number of clients on a logarithmic scale.

The legend identifies the following policies:

- linux
- cfs_wwc
- cfs_wwc_dumb_new_unblock
- cfs_wwc_flat_4msLB
- cfs_wwc_flat_4msLB_lock scheduling policy, cfs_wwc_flat_4msLB_lock scheduling policy
- cfs_wwc_flat_v2
- cfs_wwc_local_new_unblock scheduling policy, cfs_wwc_local_new_unblock scheduling policy
- cfs_wwc_lookalike
- ule
- ule_wwc_v3

The 'ule' policy (yellow line) shows a significant increase in the metric as the number of clients increases, reaching a value of approximately 95 at 512 clients. The 'ule_wwc_v3' policy (cyan line) shows a much more gradual increase, reaching a value of approximately 15 at 512 clients. All other policies (linux, cfs_wwc, cfs_wwc_dumb_new_unblock, cfs_wwc_flat_4msLB, cfs_wwc_flat_4msLB_lock scheduling policy, cfs_wwc_flat_v2, cfs_wwc_local_new_unblock scheduling policy, cfs_wwc_local_new_unblock scheduling policy, cfs_wwc_lookalike) show a very slight increase in the metric as the number of clients increases, remaining below 5 across the entire range.