

Figure 1: Number of clients vs. number of requests per client. The plot shows the performance of various scheduling policies as the number of clients increases from 32 to 512. The y-axis represents a metric (likely latency or throughput) ranging from 0 to 1000. The x-axis represents the number of clients, with major ticks at 32, 64, 128, 256, and 512. The legend identifies the following policies:

- linux (blue)
- cfs_wwc (orange)
- cfs_wwc_dumb_new_unblock (green)
- cfs_wwc_flat_4msLB (red)
- cfs_wwc_flat_4msLB_lock (purple)
- cfs_wwc_flat_v2 (brown)
- cfs_wwc_local_new_unblock (pink)
- cfs_wwc_lookalike (grey)
- ule (yellow)
- ule_wwc_v3 (teal)

The 'ule' policy (yellow line) shows a significant increase in the metric as the number of clients increases, reaching approximately 1000 at 512 clients. The 'ule_wwc_v3' policy (teal line) shows a much more gradual increase, reaching approximately 200 at 512 clients. The other policies (linux, cfs_wwc, cfs_wwc_dumb_new_unblock, cfs_wwc_flat_4msLB, cfs_wwc_flat_4msLB_lock, cfs_wwc_flat_v2, cfs_wwc_local_new_unblock, cfs_wwc_lookalike) all show a very low and relatively constant metric value, near zero, across the range of clients.

